

On the Foundations of Document Understanding

Ada Lovelace, Alan Turing

Abstract

We study the problem of extracting structured information from research papers. Our approach combines a vision-language model with a downstream agent. Experiments on a small held-out set indicate the approach is promising.

1. Introduction